

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as py
import seaborn as sns
from sklearn.linear_model import LogisticRegression
```

```
In [2]: df = pd.read_csv(r"D:\datasets\madrid_2008.csv")
df
```

```
Out[2]:
```

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY
0	2008-06-01 01:00:00	NaN	0.47	NaN	NaN	NaN	83.089996	120.699997	NaN
1	2008-06-01 01:00:00	NaN	0.59	NaN	NaN	NaN	94.820000	130.399994	NaN
2	2008-06-01 01:00:00	NaN	0.55	NaN	NaN	NaN	75.919998	104.599998	NaN
3	2008-06-01 01:00:00	NaN	0.36	NaN	NaN	NaN	61.029999	66.559998	NaN
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61
...
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91
226388	2008-11-01 00:00:00	NaN	0.30	NaN	NaN	NaN	41.880001	48.500000	NaN
226389	2008-11-01 00:00:00	0.25	NaN	0.56	NaN	0.11	83.610001	102.199997	NaN
226390	2008-11-01 00:00:00	0.54	NaN	2.70	NaN	0.18	70.639999	81.860001	NaN
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64

226392 rows × 17 columns

```
In [3]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 226392 entries, 0 to 226391
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        226392 non-null  object
1   BEN         67047 non-null   float64
2   CO          208109 non-null  float64
3   EBE         67044 non-null   float64
4   MXY         25867 non-null   float64
5   NMHC        85079 non-null   float64
6   NO_2        225315 non-null  float64
7   NOx         225311 non-null  float64
8   OXY         25878 non-null   float64
9   O_3         215716 non-null  float64
10  PM10        220179 non-null  float64
11  PM25        67833 non-null   float64
12  PXY         25877 non-null   float64
13  SO_2        225405 non-null  float64
14  TCH         85107 non-null   float64
15  TOL         66940 non-null   float64
16  station     226392 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 29.4+ MB

```

```

In [4]: df1 = df.fillna(value=0)
df1

```

Out[4]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY
0	2008-06-01 01:00:00	0.00	0.47	0.00	0.00	0.00	83.089996	120.699997	0.00
1	2008-06-01 01:00:00	0.00	0.59	0.00	0.00	0.00	94.820000	130.399994	0.00
2	2008-06-01 01:00:00	0.00	0.55	0.00	0.00	0.00	75.919998	104.599998	0.00
3	2008-06-01 01:00:00	0.00	0.36	0.00	0.00	0.00	61.029999	66.559998	0.00
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61
...
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91
226388	2008-11-01 00:00:00	0.00	0.30	0.00	0.00	0.00	41.880001	48.500000	0.00
226389	2008-11-01 00:00:00	0.25	0.00	0.56	0.00	0.11	83.610001	102.199997	0.00
226390	2008-11-01 00:00:00	0.54	0.00	2.70	0.00	0.18	70.639999	81.860001	0.00
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64

226392 rows × 17 columns

In [5]: `df1.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 226392 entries, 0 to 226391
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        226392 non-null  object
1   BEN         226392 non-null  float64
2   CO          226392 non-null  float64
3   EBE         226392 non-null  float64
4   MXY         226392 non-null  float64
5   NMHC        226392 non-null  float64
6   NO_2        226392 non-null  float64
7   NOx         226392 non-null  float64
8   OXY         226392 non-null  float64
9   O_3         226392 non-null  float64
10  PM10        226392 non-null  float64
11  PM25        226392 non-null  float64
12  PXY         226392 non-null  float64
13  SO_2        226392 non-null  float64
14  TCH         226392 non-null  float64
15  TOL         226392 non-null  float64
16  station     226392 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 29.4+ MB

```

```
In [6]: df1.columns
```

```

Out[6]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
              'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')

```

```

In [7]: df2=df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
df2

```

Out[7]:


	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3
0	0.00	0.47	0.00	0.00	0.00	83.089996	120.699997	0.00	16.990000
1	0.00	0.59	0.00	0.00	0.00	94.820000	130.399994	0.00	17.469999
2	0.00	0.55	0.00	0.00	0.00	75.919998	104.599998	0.00	13.470000
3	0.00	0.36	0.00	0.00	0.00	61.029999	66.559998	0.00	23.110001
4	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000
...
226387	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.400002
226388	0.00	0.30	0.00	0.00	0.00	41.880001	48.500000	0.00	35.830002
226389	0.25	0.00	0.56	0.00	0.11	83.610001	102.199997	0.00	14.130000
226390	0.54	0.00	2.70	0.00	0.18	70.639999	81.860001	0.00	0.000000
226391	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.910000

226392 rows × 16 columns

In [8]: `sns.pairplot(df2)`

C:\Users\HP\AppData\Local\Programs\Python\Python311\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)

Out[8]: `<seaborn.axisgrid.PairGrid at 0x20cbf1b1b10>`



```
In [10]: x=df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL']]
         y=df2['station']
```

```
In [11]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test =train_test_split(x,y,test_size=0.3)
```

linear

```
In [12]: from sklearn.linear_model import LinearRegression
```

```
In [13]: lr=LinearRegression()
         lr.fit(x_train,y_train)
```

```
Out[13]: ▼ LinearRegression
         LinearRegression()
```

```
In [14]: coeff =pd.DataFrame(lr.coef_,x.columns,columns=["Co-efficient"])
         coeff
```

Out[14]:


Co-efficient	
BEN	-3.354127
CO	-10.509262
EBE	1.727078
MXY	-1.251906
NMHC	0.752771
NO_2	-0.070672
NOx	0.024686
OXY	3.808906
O_3	-0.030009
PM10	0.012944
PM25	0.241430
PXY	9.280505
SO_2	-0.090592
TCH	0.835774
TOL	-0.398473

In [15]: `print(lr.intercept_)`

28079026.48541442

In [16]: `prediction = lr.predict(x_test)`
`py.scatter(y_test, prediction)`

Out[16]: `<matplotlib.collections.PathCollection at 0x20c87028290>`



```
In [17]: print(lr.score(x_test,y_test))
```

```
0.1402608789458767
```

```
In [18]: print(lr.score(x_train,y_train))
```

```
0.13424011207710673
```

Ridge

```
In [19]: from sklearn.linear_model import Ridge,Lasso
```

```
In [20]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
```

```
Out[20]: ▼ Ridge
         Ridge(alpha=10)
```

```
In [21]: rr.score(x_test,y_test)
```

```
Out[21]: 0.1402434037616892
```

Lasso

```
In [22]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[22]: ▼ Lasso
         Lasso(alpha=10)
```

```
In [23]: la.score(x_test,y_test)
```

```
Out[23]: 0.045338066933052756
```

elasticnet

```
In [24]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[24]: ▼ ElasticNet
         ElasticNet()
```

```
In [25]: print(en.coef_)
```

```
[-0.         -0.         0.         1.73457161  0.         -0.05936775
 -0.01442502  1.23684634 -0.03562738  0.01699235  0.3231309   1.01080816
 -0.15080874  0.31505937 -0.1074231 ]
```

```
In [26]: print(en.intercept_)
```

```
28079026.545498498
```

```
In [27]: print(en.predict(x_test))
```

```
[28079020.75269785 28079021.23305743 28079024.5408356 ...
 28079033.19802859 28079021.48640617 28079021.73647853]
```

```
In [28]: print(en.score(x_test,y_test))
```

```
0.10169399736788343
```

logistic

```
In [29]: feature_matrix =df2.iloc[:,0:15]
         target_vector=df2.iloc[:, -1]
```

```
In [30]: feature_matrix=df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY',  
                             'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df2['station']
```

```
In [31]: feature_matrix.shape
```

```
Out[31]: (226392, 15)
```

```
In [32]: target_vector.shape
```

```
Out[32]: (226392,)
```

```
In [33]: from sklearn.preprocessing import StandardScaler
```

```
In [34]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [35]: logr = LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\Users\HP\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
Out[35]: ▼ LogisticRegression  
LogisticRegression()
```

```
In [36]: observation=[[1,2,3,4,5,6,7,8,9,11,12,13,14,15,16]]
```

```
In [37]: prediction =logr.predict(observation)  
print(prediction)
```

```
[28079099]
```

```
In [38]: logr.classes_
```

```
Out[38]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,  
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,  
                28079018, 28079019, 28079021, 28079022, 28079023, 28079024,  
                28079025, 28079026, 28079027, 28079036, 28079038, 28079039,  
                28079040, 28079099], dtype=int64)
```

```
In [39]: logr.score(fs,target_vector)
```

```
Out[39]: 0.49849376303049575
```

```
In [40]: logr.predict_proba(observation)[0][0]
```

```
Out[40]: 6.191323914672894e-109
```

```
In [41]: logr.predict_proba(observation)[0][1]
```

```
Out[41]: 1.5163026503883967e-162
```

Random forest

```
In [42]: from sklearn.ensemble import RandomForestClassifier  
from sklearn.tree import plot_tree
```

```
In [43]: x=df2.drop('station',axis=1)  
y=df2['station']
```

```
In [44]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.70)
```

```
In [45]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[45]: ▼ RandomForestClassifier  
RandomForestClassifier()
```

```
In [46]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[6,7,8,9,10],  
                    'n_estimators':[11,12,13,14,15]}
```

```
In [47]: from sklearn.model_selection import GridSearchCV
```

```
In [55]: grid_search =GridSearchCV(estimator =rfc,param_grid=parameters,cv=2,scoring=  
grid_search.fit(x_train,y_train)
```

```
Out[55]: ► GridSearchCV  
► estimator: RandomForestClassifier  
    ► RandomForestClassifier
```

```
In [56]: grid_search.best_score_
```

```
Out[56]: 0.49461848194031977
```

```
In [57]: rfc_best=grid_search.best_estimator_
```

```
In [51]: py.figure(figsize=(80,50))  
plot tree(rfc_best.estimators_[5],filled=True)
```

```

Out[51]: [Text(0.5, 0.9166666666666666, 'x[11] <= 0.08\ngini = 0.961\nsamples = 4281
6\nvalue = [2644, 2671, 2672, 2699, 2626, 2591, 2707, 2609, 2685\n2523, 208
3, 2608, 2635, 2619, 2552, 2685, 2700, 2694\n2597, 2590, 2640, 2658, 2642,
2554, 2580, 2653]'),
Text(0.25, 0.75, 'x[10] <= 0.285\ngini = 0.957\nsamples = 37833\nvalue =
[2644, 2671, 2672, 36, 2626, 2591, 2707, 2609, 2685\n2523, 2083, 2608, 263
5, 2619, 2552, 2685, 2700, 79\n2597, 2590, 2640, 2658, 2642, 2554, 2580,
1]'),
Text(0.125, 0.5833333333333333, 'x[14] <= 0.055\ngini = 0.945\nsamples = 2
9886\nvalue = [64, 2671, 2672, 16, 2626, 2591, 2707, 2609, 2685\n2523, 6, 2
608, 2635, 2619, 2552, 6, 2700, 47, 2597\n13, 2640, 2658, 20, 2554, 2580,
0]'),
Text(0.0625, 0.4166666666666667, 'x[4] <= 0.005\ngini = 0.935\nsamples = 2
5055\nvalue = [64, 2671, 2672, 16, 2626, 164, 2707, 2609, 2685\n2523, 3, 26
08, 2635, 2619, 2552, 6, 52, 47, 2597\n7, 22, 2658, 20, 2554, 2580, 0]'),
Text(0.03125, 0.25, 'x[13] <= 0.59\ngini = 0.924\nsamples = 21665\nvalue =
[64, 2671, 2672, 16, 40, 31, 2707, 45, 2685, 2523\n3, 2608, 2635, 2619, 255
2, 6, 13, 47, 2597, 7, 7\n2658, 20, 2554, 2580, 0]'),
Text(0.015625, 0.0833333333333333, 'gini = 0.924\nsamples = 21639\nvalue =
[64, 2671, 2672, 16, 20, 31, 2707, 24, 2685, 2523\n3, 2608, 2635, 2619, 2
552, 6, 13, 47, 2597, 7, 7\n2658, 20, 2554, 2580, 0]'),
Text(0.046875, 0.0833333333333333, 'gini = 0.5\nsamples = 26\nvalue = [0,
0, 0, 0, 20, 0, 0, 21, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0]'),
Text(0.09375, 0.25, 'x[8] <= 49.47\ngini = 0.534\nsamples = 3390\nvalue =
[0, 0, 0, 0, 2586, 133, 0, 2564, 0, 0, 0, 0, 0\n0, 0, 0, 39, 0, 0, 0, 15,
0, 0, 0, 0, 0]'),
Text(0.078125, 0.0833333333333333, 'gini = 0.529\nsamples = 2176\nvalue =
[0, 0, 0, 0, 1844, 82, 0, 1433, 0, 0, 0, 0, 0\n0, 0, 0, 32, 0, 0, 0, 15, 0,
0, 0, 0, 0]'),
Text(0.109375, 0.0833333333333333, 'gini = 0.509\nsamples = 1214\nvalue =
[0, 0, 0, 0, 742, 51, 0, 1131, 0, 0, 0, 0, 0\n0, 0, 0, 7, 0, 0, 0, 0, 0, 0,
0, 0, 0]'),
Text(0.1875, 0.4166666666666667, 'x[9] <= 0.69\ngini = 0.667\nsamples = 48
31\nvalue = [0, 0, 0, 0, 0, 2427, 0, 0, 0, 0, 3, 0, 0, 0\n0, 0, 2648, 0, 0,
6, 2618, 0, 0, 0, 0, 0]'),
Text(0.15625, 0.25, 'x[8] <= 7.675\ngini = 0.035\nsamples = 396\nvalue =
[0, 0, 0, 0, 0, 11, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 598, 0, 0, 0, 0, 0, 0, 0,
0, 0]'),
Text(0.140625, 0.0833333333333333, 'gini = 0.408\nsamples = 8\nvalue =
[0, 0, 0, 0, 0, 10, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 4, 0, 0, 0, 0, 0, 0, 0,
0, 0]'),
Text(0.171875, 0.0833333333333333, 'gini = 0.003\nsamples = 388\nvalue =
[0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 594, 0, 0, 0, 0, 0, 0, 0,
0, 0]'),
Text(0.21875, 0.25, 'x[1] <= 0.06\ngini = 0.664\nsamples = 4435\nvalue =
[0, 0, 0, 0, 0, 2416, 0, 0, 0, 0, 3, 0, 0, 0\n0, 0, 2050, 0, 0, 6, 2618, 0,
0, 0, 0, 0]'),
Text(0.203125, 0.0833333333333333, 'gini = 0.006\nsamples = 1650\nvalue =
[0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 6, 2618, 0, 0, 0,
0, 0]'),
Text(0.234375, 0.0833333333333333, 'gini = 0.497\nsamples = 2785\nvalue =
[0, 0, 0, 0, 0, 2414, 0, 0, 0, 0, 3, 0, 0, 0\n0, 0, 2050, 0, 0, 0, 0, 0, 0,
0, 0, 0]'),
Text(0.375, 0.5833333333333333, 'x[2] <= 0.05\ngini = 0.8\nsamples = 7947

```

```

0, 2577, 0, 0, 2622, 0, 0\n1]'),
Text(0.3125, 0.4166666666666667, 'x[9] <= 10.14\ngini = 0.673\nsamples = 5
019\nvalue = [2580, 0, 0, 20, 0, 0, 0, 0, 0, 0, 13, 0, 0\n0, 0, 2679, 0, 3
2, 0, 9, 0, 0, 2622, 0, 0, 0]'),
Text(0.28125, 0.25, 'x[8] <= 29.965\ngini = 0.646\nsamples = 671\nvalue =
[297, 0, 0, 6, 0, 0, 0, 0, 0, 11, 0, 0, 0\n0, 256, 0, 1, 0, 1, 0, 0, 51
3, 0, 0, 0]'),
Text(0.265625, 0.08333333333333333, 'gini = 0.64\nsamples = 127\nvalue =
[98, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 45, 0, 0, 0, 1, 0, 0, 58, 0,
0, 0]'),
Text(0.296875, 0.08333333333333333, 'gini = 0.625\nsamples = 544\nvalue =
[199, 0, 0, 4, 0, 0, 0, 0, 0, 0, 11, 0, 0, 0\n0, 211, 0, 1, 0, 0, 0, 0, 45
5, 0, 0, 0]'),
Text(0.34375, 0.25, 'x[10] <= 19.325\ngini = 0.671\nsamples = 4348\nvalue
= [2283, 0, 0, 14, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0\n0, 2423, 0, 31, 0, 8, 0,
0, 2109, 0, 0, 0]'),
Text(0.328125, 0.08333333333333333, 'gini = 0.665\nsamples = 2903\nvalue =
[1203, 0, 0, 13, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 1815, 0, 31, 0, 7, 0, 0,
1477, 0, 0, 0]'),
Text(0.359375, 0.08333333333333333, 'gini = 0.642\nsamples = 1445\nvalue =
[1080, 0, 0, 1, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0\n0, 608, 0, 0, 0, 1, 0, 0, 63
2, 0, 0, 0]'),
Text(0.4375, 0.4166666666666667, 'x[2] <= 1.005\ngini = 0.494\nsamples = 2
928\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2064, 0, 0, 0\n0, 0, 0, 0, 0, 2
568, 0, 0, 0, 0, 0, 1]'),
Text(0.40625, 0.25, 'x[14] <= 2.06\ngini = 0.307\nsamples = 2014\nvalue =
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 597, 0, 0, 0\n0, 0, 0, 0, 0, 2559, 0, 0, 0,
0, 0, 0]'),
Text(0.390625, 0.08333333333333333, 'gini = 0.013\nsamples = 1478\nvalue =
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 15, 0, 0, 0\n0, 0, 0, 0, 0, 2299, 0, 0, 0,
0, 0, 0]'),
Text(0.421875, 0.08333333333333333, 'gini = 0.427\nsamples = 536\nvalue =
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 582, 0, 0, 0\n0, 0, 0, 0, 0, 260, 0, 0, 0,
0, 0, 0]'),
Text(0.46875, 0.25, 'x[0] <= 0.46\ngini = 0.013\nsamples = 914\nvalue =
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1467, 0, 0, 0\n0, 0, 0, 0, 0, 9, 0, 0, 0, 0,
0, 1]'),
Text(0.453125, 0.08333333333333333, 'gini = 0.48\nsamples = 9\nvalue = [0,
0, 0, 0, 0, 0, 0, 0, 0, 6, 0, 0, 0\n0, 0, 0, 0, 0, 9, 0, 0, 0, 0, 0,
0]'),
Text(0.484375, 0.08333333333333333, 'gini = 0.001\nsamples = 905\nvalue =
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1461, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 1]'),
Text(0.75, 0.75, 'x[5] <= 22.07\ngini = 0.667\nsamples = 4983\nvalue = [0,
0, 0, 2663, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 2615, 0, 0, 0, 0, 0, 0,
0, 2652]'),
Text(0.625, 0.5833333333333333, 'x[11] <= 0.835\ngini = 0.23\nsamples = 10
63\nvalue = [0, 0, 0, 75, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 1474, 0,
0, 0, 0, 0, 0, 141]'),
Text(0.5625, 0.4166666666666667, 'x[6] <= 23.315\ngini = 0.43\nsamples = 4
58\nvalue = [0, 0, 0, 73, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 541, 0, 0,
0, 0, 0, 0, 129]'),
Text(0.53125, 0.25, 'x[9] <= 11.035\ngini = 0.25\nsamples = 331\nvalue =
[0, 0, 0, 20, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 457, 0, 0, 0, 0, 0, 0,
0, 55]'),
Text(0.515625, 0.08333333333333333, 'gini = 0.376\nsamples = 185\nvalue =

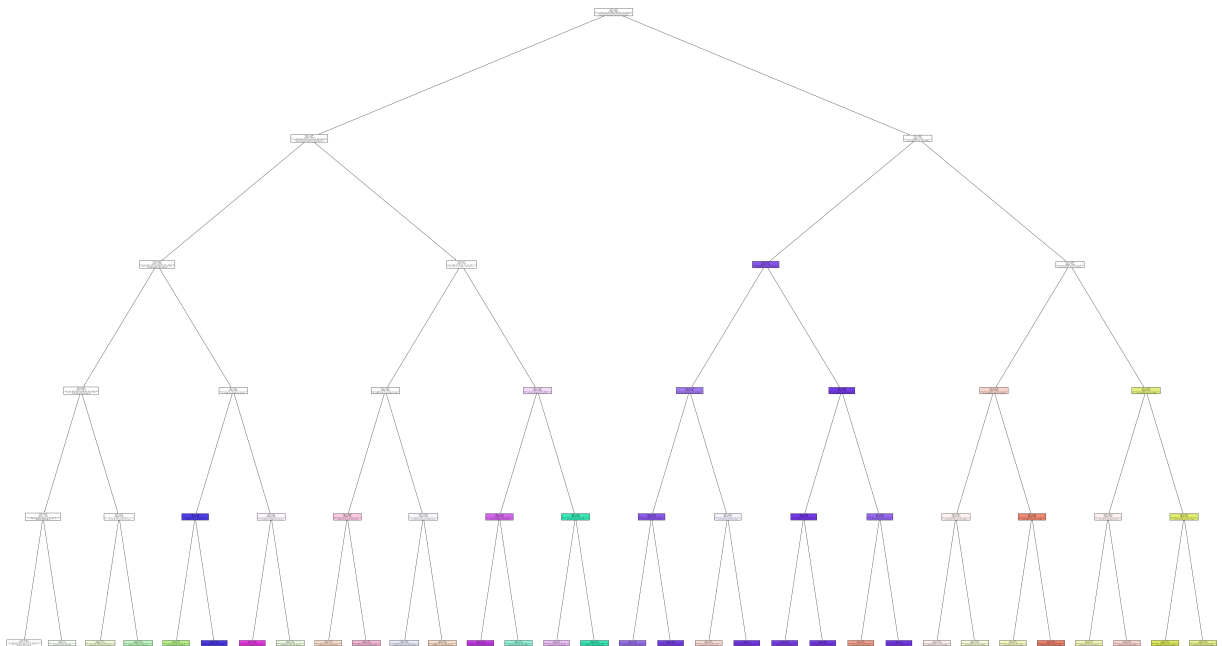
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js

```

Text(0.859375, 0.08333333333333333, 'gini = 0.365\nsamples = 623\nvalue =
[0, 0, 0, 114, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 101, 0, 0, 0, 0, 0,
0, 0, 769]'),
Text(0.9375, 0.4166666666666667, 'x[7] <= 1.435\ngini = 0.419\nsamples = 1
116\nvalue = [0, 0, 0, 1287, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 93, 0,
0, 0, 0, 0, 0, 0, 388]'),
Text(0.90625, 0.25, 'x[14] <= 7.435\ngini = 0.664\nsamples = 124\nvalue =
[0, 0, 0, 65, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 62, 0, 0, 0, 0, 0, 0,
0, 77]'),
Text(0.890625, 0.08333333333333333, 'gini = 0.587\nsamples = 42\nvalue =
[0, 0, 0, 38, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 20, 0, 0, 0, 0, 0, 0, 0,
0, 11]'),
Text(0.921875, 0.08333333333333333, 'gini = 0.624\nsamples = 82\nvalue =
[0, 0, 0, 27, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 42, 0, 0, 0, 0, 0, 0, 0,
0, 66]'),
Text(0.96875, 0.25, 'x[9] <= 41.53\ngini = 0.35\nsamples = 992\nvalue =
[0, 0, 0, 1222, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 31, 0, 0, 0, 0, 0, 0,
0, 0, 311]'),
Text(0.953125, 0.08333333333333333, 'gini = 0.247\nsamples = 532\nvalue =
[0, 0, 0, 689, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 11, 0, 0, 0, 0, 0, 0, 0,
0, 103]'),
Text(0.984375, 0.08333333333333333, 'gini = 0.434\nsamples = 460\nvalue =
[0, 0, 0, 533, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 20, 0, 0, 0, 0, 0, 0, 0,
0, 208]')

```



conclusion

The bestfit model is logistic Regression with score of 0.49849376303049575