

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: data=pd.read_csv(r"C:\Users\user\Downloads\madrid_2014.csv")
data
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2014-06-01 01:00:00	NaN	0.2	NaN	NaN	3.0	10.0	NaN	NaN	NaN	3.0	NaN	NaN	2
1	2014-06-01 01:00:00	0.2	0.2	0.1	0.11	3.0	17.0	68.0	10.0	5.0	5.0	1.36	1.3	2
2	2014-06-01 01:00:00	0.3	NaN	0.1	NaN	2.0	6.0	NaN	NaN	NaN	NaN	NaN	1.1	2
3	2014-06-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	79.0	NaN	NaN	NaN	NaN	NaN	2
4	2014-06-01 01:00:00	NaN	NaN	NaN	NaN	1.0	6.0	75.0	NaN	NaN	4.0	NaN	NaN	2
...
210019	2014-09-01 00:00:00	NaN	0.5	NaN	NaN	20.0	84.0	29.0	NaN	NaN	NaN	NaN	NaN	2
210020	2014-09-01 00:00:00	NaN	0.3	NaN	NaN	1.0	22.0	NaN	15.0	NaN	6.0	NaN	NaN	2
210021	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	1.0	13.0	70.0	NaN	NaN	NaN	NaN	NaN	2
210022	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	38.0	42.0	NaN	NaN	NaN	NaN	NaN	2
210023	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	1.0	26.0	65.0	11.0	NaN	NaN	NaN	NaN	2

210024 rows × 14 columns



```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210024 entries, 0 to 210023
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        210024 non-null  object
1   BEN         46703 non-null   float64
2   CO          87023 non-null   float64
3   EBE         46722 non-null   float64
4   NMHC        25021 non-null   float64
5   NO          209154 non-null  float64
6   NO_2        209154 non-null  float64
7   O_3         121681 non-null  float64
8   PM10        104311 non-null  float64
9   PM25        51954 non-null   float64
10  SO_2        87141 non-null   float64
11  TCH         25021 non-null   float64
12  TOL         46570 non-null   float64
13  station     210024 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: df=data.fillna(value=0)
df
```

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2014-06-01 01:00:00	0.0	0.2	0.0	0.00	3.0	10.0	0.0	0.0	0.0	3.0	0.00	0.0	28
1	2014-06-01 01:00:00	0.2	0.2	0.1	0.11	3.0	17.0	68.0	10.0	5.0	5.0	1.36	1.3	28
2	2014-06-01 01:00:00	0.3	0.0	0.1	0.00	2.0	6.0	0.0	0.0	0.0	0.0	0.00	1.1	28
3	2014-06-01 01:00:00	0.0	0.2	0.0	0.00	1.0	6.0	79.0	0.0	0.0	0.0	0.00	0.0	28
4	2014-06-01 01:00:00	0.0	0.0	0.0	0.00	1.0	6.0	75.0	0.0	0.0	4.0	0.00	0.0	28
...
210019	2014-09-01 00:00:00	0.0	0.5	0.0	0.00	20.0	84.0	29.0	0.0	0.0	0.0	0.00	0.0	28
210020	2014-09-01 00:00:00	0.0	0.3	0.0	0.00	1.0	22.0	0.0	15.0	0.0	6.0	0.00	0.0	28
210021	2014-09-01 00:00:00	0.0	0.0	0.0	0.00	1.0	13.0	70.0	0.0	0.0	0.0	0.00	0.0	28
210022	2014-09-01 00:00:00	0.0	0.0	0.0	0.00	3.0	38.0	42.0	0.0	0.0	0.0	0.00	0.0	28
210023	2014-09-01 00:00:00	0.0	0.0	0.0	0.00	1.0	26.0	65.0	11.0	0.0	0.0	0.00	0.0	28

210024 rows × 14 columns

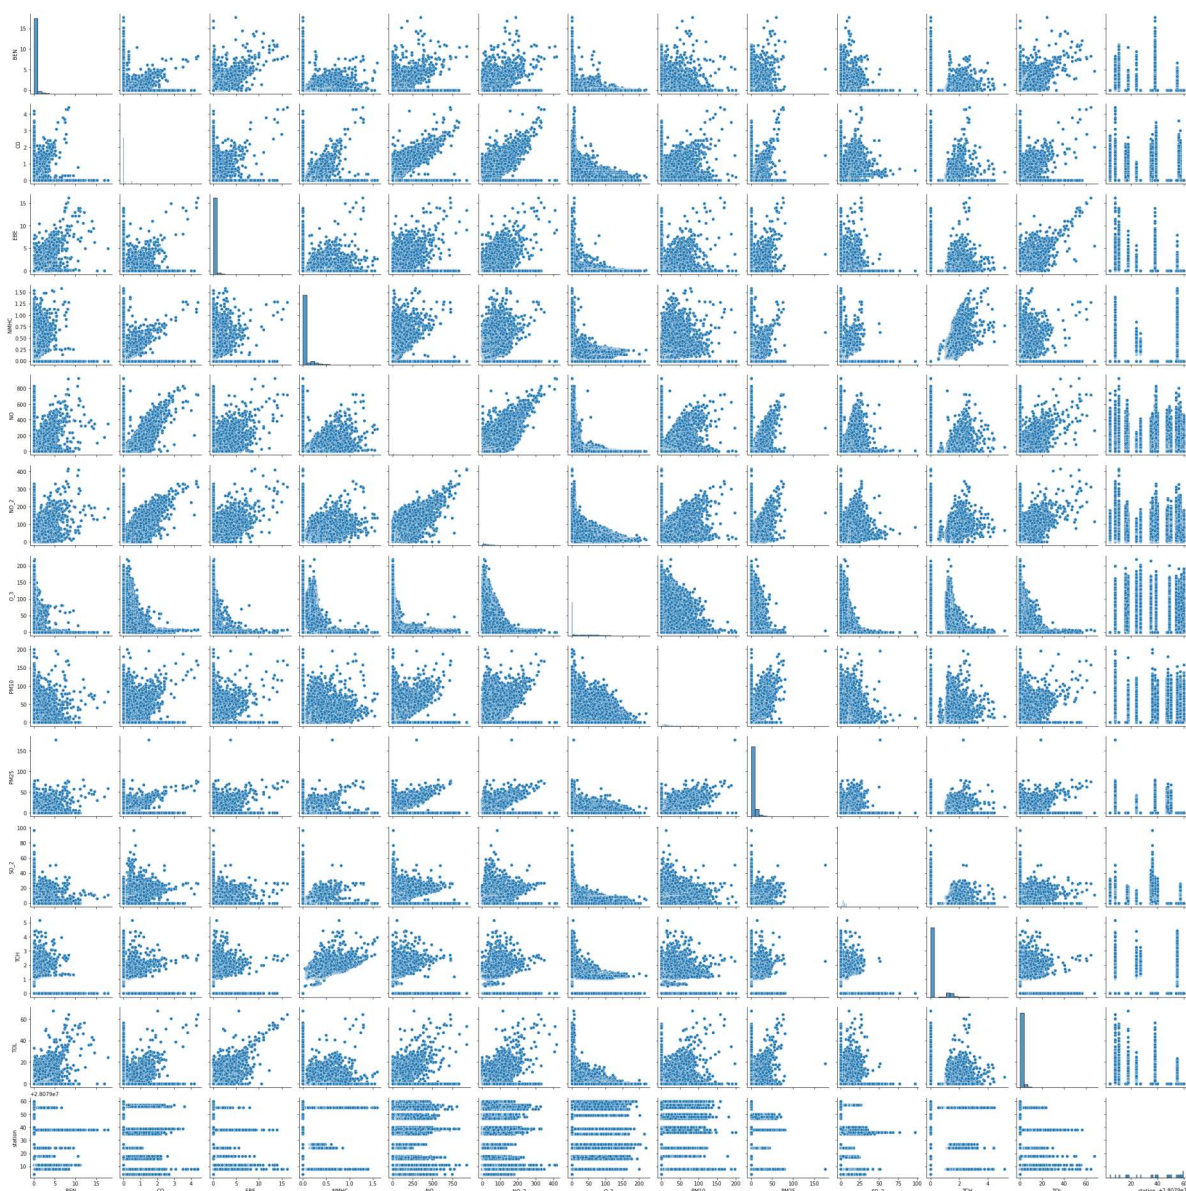


```
In [5]: df.columns
```

Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

```
In [6]: sns.pairplot(df)
```

```
Out[6]: <seaborn.axisgrid.PairGrid at 0x22900044ac0>
```




```
In [45]: sns.distplot(data["station"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[45]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



MODEL BUILDING

Linear Regression

```
In [46]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
               'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [47]: x=df1[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
               'SO_2', 'TCH', 'TOL']]  
y=df1[['station']]
```

```
In [48]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [49]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```


```
Out[49]: LinearRegression()
```

```
In [50]: print(lr.intercept_)
```

```
[28079042.00855549]
```

```
In [51]: prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[51]: <matplotlib.collections.PathCollection at 0x2290dbb2ac0>
```



```
In [52]: print(lr.score(x_test,y_test))
```

```
0.30805331568051975
```

Ridge Regression

```
In [53]: from sklearn.linear_model import Ridge
```

```
In [54]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[54]: Ridge(alpha=10)
```

```
In [55]: rr.score(x_test,y_test)
```

```
Out[55]: 0.3080823444433828
```

Lasso Regression

```
In [56]: from sklearn.linear_model import Lasso
```

```
In [57]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[57]: Lasso(alpha=10)
```

```
In [58]: la.score(x_test,y_test)
```

```
Out[58]: 0.11705608702010839
```

Elastic Regression

```
In [59]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[59]: ElasticNet()
```

```
In [60]: print(en.coef_)
```

```
[-0.68535783 -0.24077958 -0.          -0.          0.03453292 -0.02905656
 -0.01747141  0.28372955 -0.26279892 -1.7989431  -0.39956777 -1.87883152]
```

```
In [61]: print(en.predict(x_test))
```

```
[28079040.27800945 28079039.76204846 28079043.22049388 ...
 28079044.72021452 28079043.8532663  28079041.58557998]
```

```
In [62]: print(en.score(x_test,y_test))
```

```
0.2358857348730783
```

Logistic Regression

```
In [63]: from sklearn.linear_model import LogisticRegression
```

```
In [64]: feature_matrix=df1.iloc[:,0:14]
         target_vector=df1.iloc[:,-1]
```

```
In [65]: feature_matrix.shape
```

```
Out[65]: (210024, 13)
```

```
In [66]: target_vector.shape
```

```
Out[66]: (210024,)
```

```
In [67]: from sklearn.preprocessing import StandardScaler
```

```
In [68]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [69]: logr=LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[69]: LogisticRegression()
```

```
In [70]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13]]
```

```
In [71]: prediction=logr.predict(observation)  
print(observation)
```

```
[[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13]]
```

```
In [72]: logr.classes_
```

```
Out[72]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,  
                28079024, 28079027, 28079035, 28079036, 28079038, 28079039,  
                28079040, 28079047, 28079048, 28079049, 28079050, 28079054,  
                28079055, 28079056, 28079057, 28079058, 28079059, 28079060],  
              dtype=int64)
```

```
In [73]: logr.score(fs,target_vector)
```

```
Out[73]: 0.975402811107302
```

Random Forest

```
In [74]: from sklearn.ensemble import RandomForestClassifier  
from sklearn.tree import plot_tree
```



```
In [75]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
              'SO_2', 'TCH', 'TOL', 'station']]  
x=df1[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
       'SO_2', 'TCH', 'TOL']]  
y=df1['station']
```

```
In [76]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.70)
```

```
In [77]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[77]: RandomForestClassifier()

```
In [78]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]}
```

```
In [79]: from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring='acc  
grid_search.fit(x_train,y_train)
```

Out[79]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
 param_grid={'max_depth': [1, 2, 3, 4, 5],
 'min_samples_leaf': [5, 10, 15, 20, 25],
 'n_estimators': [10, 20, 30, 40, 50]},
 scoring='accuracy')

```
In [80]: grid_search.best_score_
```

Out[80]: 0.7358229767737421

```
In [81]: rfc_best=grid_search.best_estimator_
```

```
In [82]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,filled=True)

0, 0, 0, 0]'),
  Text(1913.1428571428569, 543.5999999999999, 'gini = 0.145\nsamples = 31\n
value = [0, 4, 0, 0, 0, 47, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0,
0, 0, 0]'),
  Text(2072.5714285714284, 543.5999999999999, 'TOL <= 9.0\ngini = 0.004\nsa
mples = 1569\nvalue = [0, 5, 0, 0, 0, 2495, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0,
0, 0, 0, 0, 0, 0]'),
  Text(1992.8571428571427, 181.19999999999982, 'gini = 0.002\nsamples = 152
9\nvalue = [0, 2, 0, 0, 0, 2433, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0,
0, 0, 0, 0]'),
  Text(2152.285714285714, 181.19999999999982, 'gini = 0.088\nsamples = 40\n
value = [0, 3, 0, 0, 0, 62, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 0, 0, 0,
0, 0, 0]'),
  Text(3407.785714285714, 1630.8000000000002, 'BEN <= 0.05\ngini = 0.833\ns
amples = 9881\nvalue = [0, 2633, 0, 0, 0, 0, 2517, 0, 0, 0, 2594, 0, 0\n25
62, 2700, 0, 2589, 0, 0, 0, 0, 0, 0, 0]'),
  Text(2869.7142857142853, 1268.4, 'PM25 <= 3.5\ngini = 0.694\nsamples = 52
06\nvalue = [0, 78, 0, 0, 0, 0, 153, 0, 0, 0, 124, 0, 0\n2562, 2700, 0, 25
89, 0, 0, 0, 0, 0, 0, 0]'),
  Text(2550.8571428571427, 906.0, 'NO 2 <= 10.5\ngini = 0.545\nsamples = 55
```

Results

The best model is Logistic Regression 0.975402811107302

In []: