


```
In [64]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [65]: data=pd.read_csv(r"C:\Users\user\Downloads\madrid_2015.csv")
data
```

```
Out[65]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2015-10-01 01:00:00	NaN	0.8	NaN	NaN	90.0	82.0	NaN	NaN	NaN	10.0	NaN	NaN	2
1	2015-10-01 01:00:00	2.0	0.8	1.6	0.33	40.0	95.0	4.0	37.0	24.0	12.0	1.83	8.3	2
2	2015-10-01 01:00:00	3.1	NaN	1.8	NaN	29.0	97.0	NaN	NaN	NaN	NaN	NaN	7.1	2
3	2015-10-01 01:00:00	NaN	0.6	NaN	NaN	30.0	103.0	2.0	NaN	NaN	NaN	NaN	NaN	2
4	2015-10-01 01:00:00	NaN	NaN	NaN	NaN	95.0	96.0	2.0	NaN	NaN	9.0	NaN	NaN	2
...
210091	2015-08-01 00:00:00	NaN	0.2	NaN	NaN	11.0	33.0	53.0	NaN	NaN	NaN	NaN	NaN	2
210092	2015-08-01 00:00:00	NaN	0.2	NaN	NaN	1.0	5.0	NaN	26.0	NaN	10.0	NaN	NaN	2
210093	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	1.0	7.0	74.0	NaN	NaN	NaN	NaN	NaN	2
210094	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	3.0	7.0	65.0	NaN	NaN	NaN	NaN	NaN	2
210095	2015-08-01 00:00:00	NaN	NaN	NaN	NaN	1.0	9.0	54.0	29.0	NaN	NaN	NaN	NaN	2

210096 rows × 14 columns



```
In [66]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210096 entries, 0 to 210095
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        210096 non-null  object
1   BEN         51039 non-null   float64
2   CO          86827 non-null   float64
3   EBE         50962 non-null   float64
4   NMHC        25756 non-null   float64
5   NO          208805 non-null   float64
6   NO_2        208805 non-null   float64
7   O_3         121574 non-null   float64
8   PM10        102745 non-null   float64
9   PM25        48798 non-null   float64
10  SO_2        86898 non-null   float64
11  TCH         25756 non-null   float64
12  TOL         50626 non-null   float64
13  station     210096 non-null   int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [67]: df=data.fillna(value=0)
df
```

```
Out[67]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2015-10-01 01:00:00	0.0	0.8	0.0	0.00	90.0	82.0	0.0	0.0	0.0	10.0	0.00	0.0	28
1	2015-10-01 01:00:00	2.0	0.8	1.6	0.33	40.0	95.0	4.0	37.0	24.0	12.0	1.83	8.3	28
2	2015-10-01 01:00:00	3.1	0.0	1.8	0.00	29.0	97.0	0.0	0.0	0.0	0.0	0.00	7.1	28
3	2015-10-01 01:00:00	0.0	0.6	0.0	0.00	30.0	103.0	2.0	0.0	0.0	0.0	0.00	0.0	28
4	2015-10-01 01:00:00	0.0	0.0	0.0	0.00	95.0	96.0	2.0	0.0	0.0	9.0	0.00	0.0	28
...
210091	2015-08-01 00:00:00	0.0	0.2	0.0	0.00	11.0	33.0	53.0	0.0	0.0	0.0	0.00	0.0	28
210092	2015-08-01 00:00:00	0.0	0.2	0.0	0.00	1.0	5.0	0.0	26.0	0.0	10.0	0.00	0.0	28
210093	2015-08-01 00:00:00	0.0	0.0	0.0	0.00	1.0	7.0	74.0	0.0	0.0	0.0	0.00	0.0	28
210094	2015-08-01 00:00:00	0.0	0.0	0.0	0.00	3.0	7.0	65.0	0.0	0.0	0.0	0.00	0.0	28
210095	2015-08-01 00:00:00	0.0	0.0	0.0	0.00	1.0	9.0	54.0	29.0	0.0	0.0	0.00	0.0	28

210096 rows × 14 columns

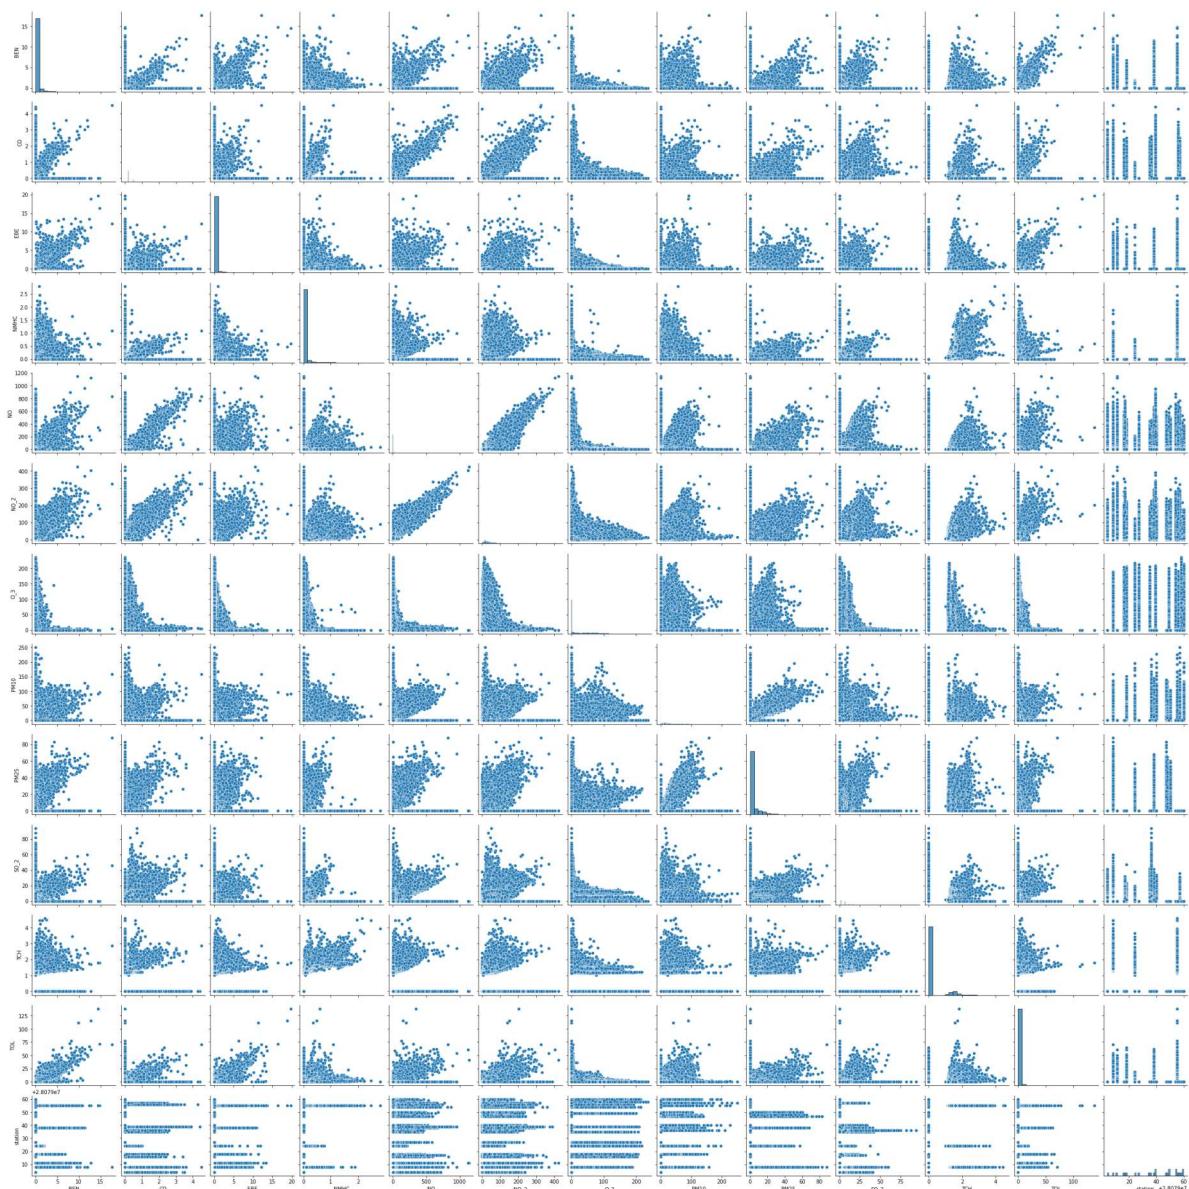


```
In [68]: df.columns
```

```
Out[68]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
               'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [69]: sns.pairplot(df)
```

```
Out[69]: <seaborn.axisgrid.PairGrid at 0x1f885d45f40>
```



```
In [ ]: sns.distplot(data["station"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

MODEL BUILDING

Linear Regression

```
In [57]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
               'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [58]: x=df1[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
               'SO_2', 'TCH', 'TOL']]
y=df1['station']
```

```
In [59]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [60]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```


Out[60]: LinearRegression()

```
In [61]: print(lr.intercept_)
```

28079043.051533587

```
In [62]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[62]: <matplotlib.collections.PathCollection at 0x1f8855bdac0>



```
In [63]: print(lr.score(x_test,y_test))
```

0.29815273412055265

Ridge Regression

```
In [17]: from sklearn.linear_model import Ridge
```

```
In [18]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]: rr.score(x_test,y_test)
```

```
Out[19]: 0.2990637039599836
```

Lasso Regression

```
In [20]: from sklearn.linear_model import Lasso
```

```
In [21]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[21]: Lasso(alpha=10)
```

```
In [22]: la.score(x_test,y_test)
```

```
Out[22]: 0.14425099333511227
```

Elastic Regression

```
In [23]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[23]: ElasticNet()
```

```
In [24]: print(en.coef_)
```

```
[-6.55959003e-01 -4.26230045e-01 -7.95802910e-01 -0.00000000e+00
 2.91478073e-02 -3.41943221e-02 -1.14131981e-03 2.78342122e-01
-2.46922131e-01 -1.31701512e+00 -0.00000000e+00 -1.69251706e+00]
```

```
In [25]: print(en.predict(x_test))
```

```
[28079035.77703039 28079035.11076136 28079040.3825461 ...
28079032.62400804 28079038.38271381 28079041.36218826]
```

```
In [26]: print(en.score(x_test,y_test))
```

```
0.2418437372373553
```

Logistic Regression

```
In [27]: from sklearn.linear_model import LogisticRegression
```

```
In [28]: feature_matrix=df1.iloc[:,0:14]
target_vector=df1.iloc[:,-1]
```

```
In [29]: feature_matrix.shape
```

```
Out[29]: (209928, 13)
```

```
In [30]: target_vector.shape
```

```
Out[30]: (209928,)
```

```
In [31]: from sklearn.preprocessing import StandardScaler
```

```
In [32]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [33]: logr=LogisticRegression()
logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[33]: LogisticRegression()
```

```
In [36]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13]]
```

```
In [37]: prediction=logr.predict(observation)
print(observation)
```

```
[[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13]]
```

```
In [38]: logr.classes_
```

```
Out[38]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,
                28079024, 28079027, 28079035, 28079036, 28079038, 28079039,
                28079040, 28079047, 28079048, 28079049, 28079050, 28079054,
                28079055, 28079056, 28079057, 28079058, 28079059, 28079060],
                dtype=int64)
```

```
In [39]: logr.score(fs,target_vector)
```

```
Out[39]: 0.9833323806257384
```

Random Forest

```
In [40]: from sklearn.ensemble import RandomForestClassifier
         from sklearn.tree import plot_tree
```

```
In [41]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
                'SO_2', 'TCH', 'TOL', 'station']]
         x=df1[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
                'SO_2', 'TCH', 'TOL']]
         y=df1['station']
```

```
In [42]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.70)
```

```
In [43]: rfc=RandomForestClassifier()
         rfc.fit(x_train,y_train)
```

```
Out[43]: RandomForestClassifier()
```

```
In [44]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
```

```
In [45]: from sklearn.model_selection import GridSearchCV
         grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring='acc
         grid_search.fit(x_train,y_train)
```

```
Out[45]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [46]: grid_search.best_score_
```

```
Out[46]: 0.7565975419987933
```



```
In [47]: rfc_best=grid_search.best_estimator_
```

```
In [48]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,filled=True)
```

```
Out[48]: [Text(2222.7000000000003, 1993.2, 'SO_2 <= 0.5\ngini = 0.958\nsamples = 39
927\nvalue = [2502, 2640, 2607, 2536, 2708, 2675, 2644, 2669, 2692\n2559,
2474, 2484, 2677, 2724, 2550, 2674, 2714, 2597\n2579, 2756, 2668, 2687, 25
12, 2650]'),
Text(1171.8000000000002, 1630.8000000000002, 'O_3 <= 0.5\ngini = 0.929\ns
amples = 23321\nvalue = [8, 16, 2607, 2536, 10, 3, 16, 2669, 13, 2, 17\n24
84, 14, 2724, 2550, 2674, 2714, 2597, 2579, 2756\n1, 2687, 2512, 2650]'),
Text(595.2, 1268.4, 'BEN <= 0.05\ngini = 0.805\nsamples = 8427\nvalue =
[8, 4, 2607, 4, 8, 0, 13, 4, 2, 2, 17, 2, 14\n2724, 2550, 8, 2714, 5, 257
9, 20, 1, 15, 27, 3]'),
Text(297.6, 906.0, 'TCH <= 0.625\ngini = 0.685\nsamples = 5184\nvalue =
[8, 1, 54, 4, 8, 0, 8, 4, 2, 2, 16, 2, 14\n2724, 2550, 8, 2714, 5, 30, 20,
1, 15, 27, 3]'),
Text(148.8, 543.5999999999999, 'NO <= 5.5\ngini = 0.682\nsamples = 5156\n
value = [8, 1, 54, 4, 8, 0, 8, 1, 2, 2, 16, 2, 14\n2724, 2550, 8, 2714, 5,
0, 20, 1, 12, 27, 3]'),
Text(74.4, 181.19999999999982, 'gini = 0.666\nsamples = 2304\nvalue = [5,
1, 31, 1, 5, 0, 8, 0, 1, 0, 16, 1, 13\n1550, 1198, 8, 728, 5, 0, 17, 0, 8,
20, 1]'),
Text(1633.0000000000002, 1631.0000000000002, 'SO_2 <= 0.657\nsamples = 335
927\nvalue = [2502, 2640, 2607, 2536, 2708, 2675, 2644, 2669, 2692\n2559,
2474, 2484, 2677, 2724, 2550, 2674, 2714, 2597\n2579, 2756, 2668, 2687, 25
12, 2650]')]
```

Results

The best model is Logistic Regression 0.9833323806257384

```
In [ ]:
```