

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: data=pd.read_csv(r"C:\Users\user\Downloads\madrid_2016.csv")
data
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL
0	2016-11-01 01:00:00	NaN	0.7	NaN	NaN	153.0	77.0	NaN	NaN	NaN	7.0	NaN	NaN
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4
2	2016-11-01 01:00:00	5.9	NaN	7.5	NaN	297.0	139.0	NaN	NaN	NaN	NaN	NaN	26.0
3	2016-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	113.0	2.0	NaN	NaN	NaN	NaN	NaN
4	2016-11-01 01:00:00	NaN	NaN	NaN	NaN	275.0	127.0	2.0	NaN	NaN	18.0	NaN	NaN
...
209491	2016-07-01 00:00:00	NaN	0.2	NaN	NaN	2.0	29.0	73.0	NaN	NaN	NaN	NaN	NaN
209492	2016-07-01 00:00:00	NaN	0.3	NaN	NaN	1.0	29.0	NaN	36.0	NaN	5.0	NaN	NaN
209493	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	1.0	19.0	71.0	NaN	NaN	NaN	NaN	NaN
209494	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	6.0	17.0	85.0	NaN	NaN	NaN	NaN	NaN
209495	2016-07-01 00:00:00	NaN	NaN	NaN	NaN	2.0	46.0	61.0	34.0	NaN	NaN	NaN	NaN

209496 rows × 14 columns



```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209496 entries, 0 to 209495
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209496 non-null  object
1   BEN         50755 non-null   float64
2   CO          85999 non-null   float64
3   EBE         50335 non-null   float64
4   NMHC        25970 non-null   float64
5   NO          208614 non-null  float64
6   NO_2        208614 non-null  float64
7   O_3         121197 non-null  float64
8   PM10        102892 non-null  float64
9   PM25        52165 non-null   float64
10  SO_2        86023 non-null   float64
11  TCH         25970 non-null   float64
12  TOL         50662 non-null   float64
13  station     209496 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: df=data.fillna(value=0)
df
```

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2016-11-01 01:00:00	0.0	0.7	0.0	0.00	153.0	77.0	0.0	0.0	0.0	7.0	0.00	0.0	2
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	2
2	2016-11-01 01:00:00	5.9	0.0	7.5	0.00	297.0	139.0	0.0	0.0	0.0	0.0	0.00	26.0	2
3	2016-11-01 01:00:00	0.0	1.0	0.0	0.00	154.0	113.0	2.0	0.0	0.0	0.0	0.00	0.0	2
4	2016-11-01 01:00:00	0.0	0.0	0.0	0.00	275.0	127.0	2.0	0.0	0.0	18.0	0.00	0.0	2
...
209491	2016-07-01 00:00:00	0.0	0.2	0.0	0.00	2.0	29.0	73.0	0.0	0.0	0.0	0.00	0.0	2
209492	2016-07-01 00:00:00	0.0	0.3	0.0	0.00	1.0	29.0	0.0	36.0	0.0	5.0	0.00	0.0	2
209493	2016-07-01 00:00:00	0.0	0.0	0.0	0.00	1.0	19.0	71.0	0.0	0.0	0.0	0.00	0.0	2
209494	2016-07-01 00:00:00	0.0	0.0	0.0	0.00	6.0	17.0	85.0	0.0	0.0	0.0	0.00	0.0	2
209495	2016-07-01 00:00:00	0.0	0.0	0.0	0.00	2.0	46.0	61.0	34.0	0.0	0.0	0.00	0.0	2

209496 rows × 14 columns

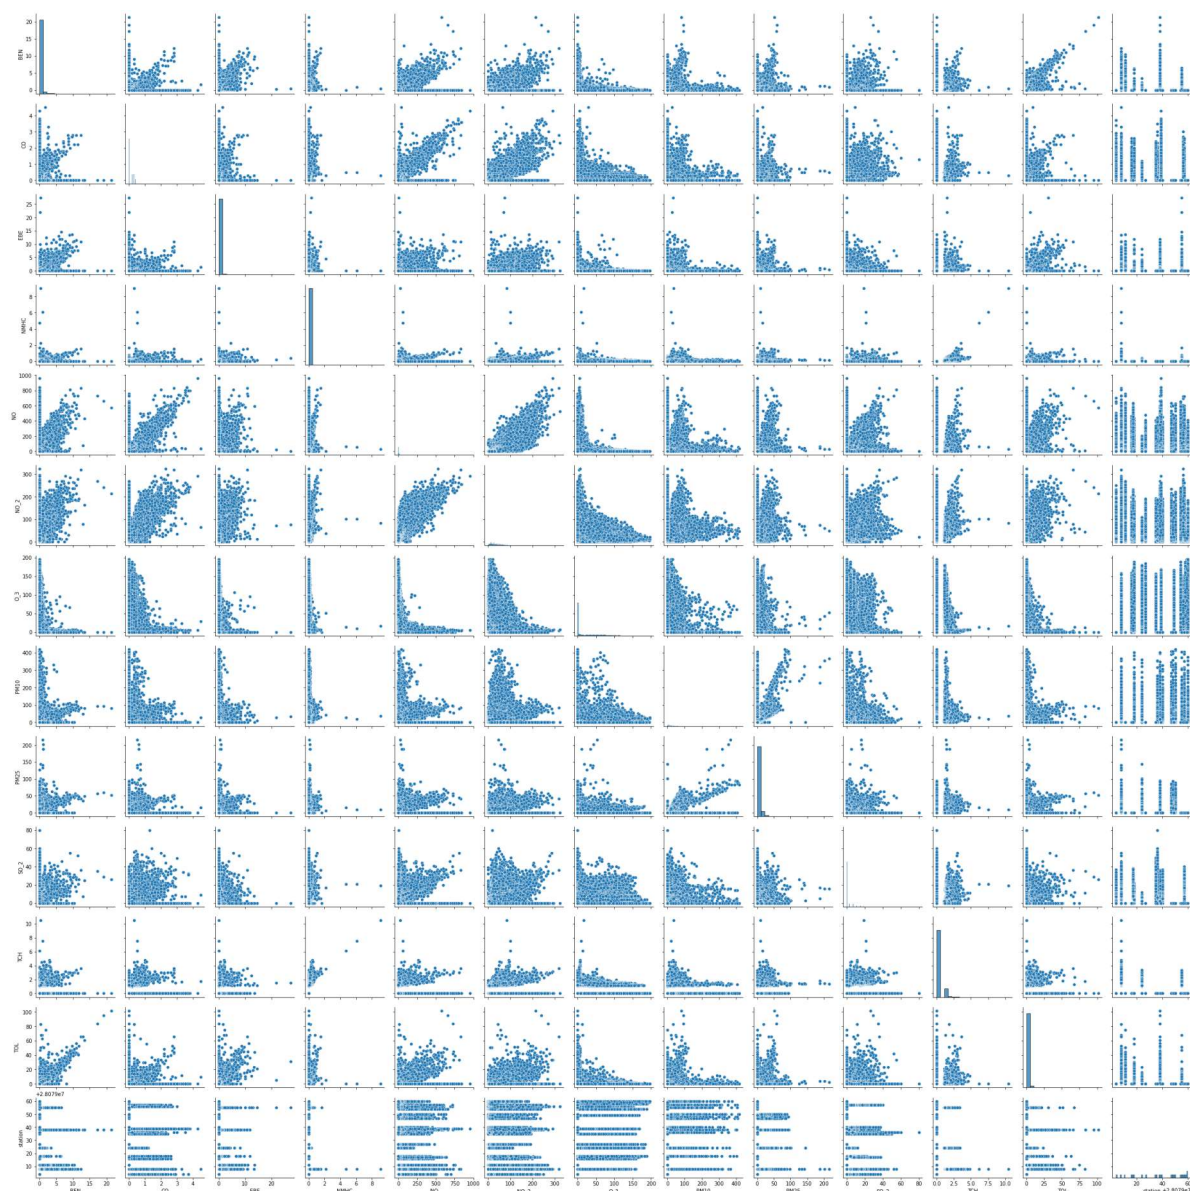


```
In [5]: df.columns
```

Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

```
In [6]: sns.pairplot(df)
```

```
Out[6]: <seaborn.axisgrid.PairGrid at 0x18223fe4730>
```




```
In [8]: sns.distplot(data["station"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[8]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



MODEL BUILDING

Linear Regression

```
In [10]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
               'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [11]: x=df1[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
               'SO_2', 'TCH', 'TOL']]  
y=df1[['station']]
```

```
In [12]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [13]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```


```
Out[13]: LinearRegression()
```

```
In [14]: print(lr.intercept_)
```

```
[28079043.29340011]
```

```
In [15]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x18233e09a00>
```



```
In [16]: print(lr.score(x_test,y_test))
```

```
0.32782218723180634
```

Ridge Regression

```
In [17]: from sklearn.linear_model import Ridge
```

```
In [18]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]: rr.score(x_test,y_test)
```

```
Out[19]: 0.3273460296426014
```

Lasso Regression

```
In [20]: from sklearn.linear_model import Lasso
```

```
In [21]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[21]: Lasso(alpha=10)
```

```
In [22]: la.score(x_test,y_test)
```

```
Out[22]: 0.16202886515402581
```

Elastic Regression

```
In [23]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[23]: ElasticNet()
```

```
In [24]: print(en.coef_)
```

```
[-0.91880846 -0.35807898 -0.          0.          0.036925 -0.04430166
 -0.0021193   0.21257535 -0.26272722 -1.30830156 -0.55258298 -1.39872323]
```

```
In [25]: print(en.predict(x_test))
```

```
[28079018.92979379 28079042.72678536 28079041.68022207 ...
 28079032.68586449 28079031.16860852 28079035.06668555]
```

```
In [26]: print(en.score(x_test,y_test))
```

```
0.24487292196921295
```

Logistic Regression

```
In [27]: from sklearn.linear_model import LogisticRegression
```

```
In [28]: feature_matrix=df1.iloc[:,0:15]
         target_vector=df1.iloc[:,-1]
```

```
In [29]: feature_matrix.shape
```

```
Out[29]: (209496, 13)
```

```
In [30]: target_vector.shape
```

```
Out[30]: (209496,)
```

```
In [31]: from sklearn.preprocessing import StandardScaler
```

```
In [32]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [33]: logr=LogisticRegression()  
logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[33]: LogisticRegression()
```

```
In [36]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13]]
```

```
In [37]: prediction=logr.predict(observation)  
print(observation)
```

```
[[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13]]
```

```
In [38]: logr.classes_
```

```
Out[38]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,  
                28079024, 28079027, 28079035, 28079036, 28079038, 28079039,  
                28079040, 28079047, 28079048, 28079049, 28079050, 28079054,  
                28079055, 28079056, 28079057, 28079058, 28079059, 28079060],  
              dtype=int64)
```

```
In [39]: logr.score(fs,target_vector)
```

```
Out[39]: 0.9837944399893077
```

Random Forest

```
In [40]: from sklearn.ensemble import RandomForestClassifier  
from sklearn.tree import plot_tree
```



```
In [44]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
              'SO_2', 'TCH', 'TOL', 'station']]  
x=df1[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
       'SO_2', 'TCH', 'TOL', 'station']]  
y=df1['station']
```

```
In [45]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.70)
```

```
In [46]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[46]: RandomForestClassifier()

```
In [47]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]}
```

```
In [48]: from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring='acc  
grid_search.fit(x_train,y_train)
```

Out[48]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
 param_grid={'max_depth': [1, 2, 3, 4, 5],
 'min_samples_leaf': [5, 10, 15, 20, 25],
 'n_estimators': [10, 20, 30, 40, 50]},
 scoring='accuracy')

```
In [49]: grid_search.best_score_
```

Out[49]: 0.9550661914460286

```
In [50]: rfc_best=grid_search.best_estimator_
```

[illegible]

Results

The best model is Logistic Regression 0.9837944399893077

In []: